

A Federated Approach to Identifying Advanced Persistent Security Threats on Enterprise Computer Networks

PROJECT PRESENTATION

Wade Wesolowsky
October 25, 2019

Acknowledgements

The Graduate Student Research Fund at Athabasca University for providing financial support to acquire computer hardware and networking components.

This research was not sponsored by any free or open source projects mentioned herein, and use of their representative symbols does not imply endorsement of any kind.

Motivation

Personal desire to learn about Computer Security

Constant media coverage of computer malware, specifically Advanced Persistent Threats

For the widest applicability of the research, it was imperative to investigate Enterprise (Windows) Computer Networks

Purpose of the Research

- Study of existing tools and methods which can be leveraged to detect (and remove) Advanced Persistent Security Threats (APST) within enterprise computer networks
- Implement a select group of tools in a federated manner to mitigate security threats and vulnerabilities
- Produce a working test system and test on a virtualized network
- Use simulated threats to see how system reports/reacts to them

Scope of Research

- Develop a way for the tools to share their intelligence
- Use mainly free and open source tools
- Simulate an enterprise (Windows) corporate environment
- Use APST simulators to test tools



Nature of Research

- The research was very open-ended and unrestrained
- This is an active research area, in which new concepts are being continually developed
 - “Threat Hunting”
- Most (computer and software) systems we are dealing with are in a *constant* state of flux



Anticipated Contributions

- Produce an amalgamation of free and open source tools which can be used for detection
 - Modify tools to build our solution
- Provide setup documentation
- Tool source code
- Making it all publicly accessible



Literature Review

- APST Lifecycle has not been finalized, so we looked at various life cycles to find one we liked
- We found many high level recommendations of how to detect APST but very little implementations of detecting APST
- Many rely on examining/monitoring the network itself
- Non-technical solutions
 - User education

Literature Review

- Most corporate networks are already compromised
 - Assume you are already compromised mindset
- Created a “mind map” to show interconnected ideas
- Use of honeypots to act as high value targets

Literature Review

- Open Source Security Information Management (OSSIM)
 - Not specifically designed with APST in mind
- Very little research exists for removal
 - Re-image / rebuild compromised systems
 - Restore from backups
 - “Process” for dealing with compromises

Gaps to Fill

- Most software solutions are commercial software or are open source software sold as part of commercial packages
 - Costly
 - Source code not available
- Counteracting malware in general, not APST specifically
- A single solution is not enough!
- Need a historical record to examine, when something goes wrong
 - Logging system is ideal for audit trail




Advanced Persistent [Security] Threats

- Advanced
 - Sophisticated, multi-stage nature
- Persistent
 - Establish non-volatile, permanent foothold in network
 - Long lasting, slow spread
- Threat
 - Sneaky / undetected exploits of the computer systems



Some Present Conundrums

- Wired networks are becoming “obsolete”
 - Network medium moving to wireless
 - Wireless transmission is “invisible” to our senses
- Valuable data is a target
 - Changing internal data
- New vulnerabilities are being discovered
 - Attacks on BIOS and firmware
- System complexity
- Aging infrastructure
- Malware countermeasures
- Cyberwarfare



APST Lifecycle Models

- Various models exist [1] and [2]
 - [1] shows various models and abstracts a recommended model
- We choose to go with [2] for simplicity
- Consists of four stages:
 - Prepare Stage
 - Access Stage
 - Resident Stage
 - Harvest Stage




Prepare Stage

- Collect intelligence on the target
 - Open source intelligence
 - <https://www.google.com>
 - <https://shodan.io>
 - Port scanning, etc.
- Investigate ways for initial compromise



Access Stage

- Initial compromise of the fortress perimeter
 - Whaling (C-level Spear phishing)
 - Waterhole Attack (Malicious, frequently visited website)
 - “Lost” USBs
 - Insider threats
 - Large attack surface
 - Default passwords
 - Complex infrastructure
- Using exploits
 - Public vulnerabilities
 - Zero-days (Undisclosed vulnerabilities)



Resident Stage

- Find places to hide in the network
- Bring in more exploit and remote access tools
- Explore the network, expand the web
- Exploit, decide what to steal...



Harvest Stage

- Exfiltrate the valuable data
 - “Phone Home”
- Obfuscation of residual evidence
 - Destroy the evidence of infiltration



Search for Detection Tools

- Need a central way to understand the collected intelligence
- We considered Splunk, but the offering was not free
- ELK / Elastic Stack meets the requirements:
 - Logstash gathers the logs and stores them centrally
 - Elasticsearch indexes the logs and makes them searchable
 - Kibana offers a graphical front-end for tables, graphs, charts, etc.
- Together they are called “ELK” for Elasticsearch, Logstash, and Kibana

Some Desirable Traits for Selected Tools

- Open Source or Free Software
 - Our project was focused on using free and open source software
- Proven Technology
 - Has the solution proven itself in the “real world”
- Active Developer Community
 - Can we get support for questions?
 - Is the software still being developed and refined?
- Fit for Purpose
 - Does it log to the ELK stack?
 - Will it fit with overall design?
- Cost
 - Minimal cost

pfSense (Perimeter Firewall)



- Billed as one of the best Open Source Firewalls
- Under active development with a stable and well-defined following
- Detection of Preparatory and Access stages

Snort (runs on pfSense)



- Suricata was a contender
- Snort was a match for our hardware (single processor machine with dual NICs)
- Has been around for a long time
- Detection in Preparation, Resident, Harvest stages

Sweet Security (Zeek)

- Ties together Raspberry Pi and Zeek Network Security Monitor (Formerly Bro)
- Allowed us to see what is being transmitted on the network
- A switch with port mirroring helped
- Detection in Access, Resident, Harvest stages




Wazuh (OSSEC fork)

- A open source project built on ELK and and OSSEC.
 - OSSEC client has been upgraded with better functionality, support for more OSes, etc.
- Has “alert” rules which detect predefined signs of compromise or attack
- Does host level monitoring using clients on each monitored OS
 - Data sent back to the ELK stack
- Detection in Resident, Harvest stages

Sysmon

- Provides the ability to monitor system events on Windows
 - E.g. processes, network connections, registry key changes
- Winlogbeat and/or Wazuh can move these logs over to the ELK stack for further analysis
- Detection in Resident stage



HoneyTrap

- Pretends to be a high value target on our network
- You define a configuration file to control the behaviour
- Logs can be sent to the ELK stack
- Detection in Resident stage



Windows Defender


- Unable to find a free/open source virus scanner of sufficient quality
 - ClamAV
- Windows offers a virus scanner which is “free” to use on their OS
- A virus scanner is still an important component of the overall solution
- Assists with detection during the resident phase

Alternative Tools

- We would like to mention three other suites discovered during our research which have similar or overlapping functionality to our solution
 - Elastic SIEM - <https://www.elastic.co/blog/introducing-elastic-siem> (June 2019)
 - HELK - <https://github.com/Cyb3rWard0g/HELK> (Dec 2018)
 - AlienVault OSSIM - <https://www.alienvault.com/products/ossim> (~2003)


Tools Were Configured to Work Together

- Complex process to get the tools to log and act together
- Our paper includes a brief overview of the steps
 - See references in paper for more information



Indexes

- Log data is stored in Elasticsearch indexes
- Several indexes may exist for one tool
- Allow data to be queried to find results



Indexes

Elasticsearch Index	Source Tool	Type of Logs
honeytrap*	HoneyTrap	Status Messages, Honeypot Connection Attempts
logstash-*	Zeek	Zeek Network Traffic Monitoring
pfsense-*	pfSense Firewall, Snort	Firewall Status Messages, Snort Alerts
sweet_security	Sweet Security	Detected Device Information, Port Scans
sweet_secrutiy_alerts	Sweet Security	New (Unique) Sweet Security Log Events
tardis	Sweet Security	Historial Hosts, IP Addresses, and Websites
wazuh-alerts-3.x-*	Wazuh	Log Events Above the Alert Threshold
wazuh-monitoring-3.x-*	Wazuh	All Wazuh Monitoring Logs
winlogbeats-*	Windows Event Logs	Specific Windows Event Logs (Application, System , Security, Sysmon)


Federation

- Off-the Shelf Tools (COTS) were adapted to be part of our solution
- They were configured to work together and log to the ELK stack
- Offer overlapping functionality in the various stages of APST attack

Testing

- Spend time connected to research network doing daily computing
- Simulated Enterprise (Windows) Active Directory Domain Environment
- Use tools to simulate APST attacks
 - APTSimulator (Host Level)
 - FlightSim (Network Level)
- Use tests to see examples of data returned by the tools
- Pinpoint the strengths of each tool

Testing




Querying the Data from Kibana (Plugin)

- Shows how some prebuilt queries can return data in Kibana
- Queries are predefined and return results in accessible table format
- Programmed in such a way as to be easily extensible




FSM Plugin Architecture



Querying the Data from Command-line (APT Hunter)

- Lower-level tool which is easier to debug and program than the dashboard
- Python tool to query Elasticsearch indexes
- A general query tool which can send queries to the various indexes
- Results returned give an indication of APST by counting query results



Example Detection Algorithm

Algorithm 1: Naive Detection Algorithm

Input: array *Results* of JSON Results from Elasticsearch Queries and integer of *sensors* used

Result: Prints APST Detection Results

```
1 if Results.length > 0 then
2   hits ← 0;
3   for Result ∈ Results do
4     hits = hits + Result.hits.total
5   Hits_Per_Sensor ← hits ÷ sensors;
6   APST Query Hits: hits;
7   Hits Per Sensor: Hits_Per_Sensor;
8 else
9   No APST detected;
```

Federation Continued

- APT Hunter offers the flexibility to query any/all indexes as needed
- Additional detection algorithms can be added, and then command-line arguments can be used to select between them
- Allows the tools to work together as one cohesive unit

Findings

- We have created a free and open source APST detection framework using existing security tools!
- GitHub code repository offers the ability for anyone to build on the work
- Our simulated enterprise network closely resembles a real-world enterprise network
 - Gather lessons learned to expand solution to the real world
- FSM Plugin and APT Hunter are a good starting point for the hunt

Findings

- Updating to the latest tools would help us
 - More features being added all the time to sensors
- Sysmon is very effective on the Windows host systems
 - Helps us track back the chain of events
- Using existing detection technologies would help expand the scope of investigation
 - Yara

Future Work

- Add more sensors to our network
 - SSL/TLS Decryptors
- Automation when updating the dashboard
- Machine learning to replace security analyst
- Active response
 - Alerting (SMS)



Public Dissemination

Public Code Repository on GitHub

<https://github.com/rndrev/FederatedSecurityModule>

This repository includes code used in project, project deliverables, and forked projects which we used or reference during the research.

Questions?

- What future directions do you see APST going?
- What was the most difficult part of the research?

References

- [1] B. I. D. Messaoud, K. Guennoun, M. Wahbi, and M. Sadik, “Advanced Persistent Threat: New analysis driven by life cycle phases and their challenges,” in 2016 International Conference on Advanced Communication Systems and Information Security (ACOSIS), Oct 2016, pp. 1–6. [Online]. Available: <https://doi.org/10.1109/ACOSIS.2016.7843932>
- [2] M. Li, W. Huang, Y. Wang, W. Fan, and J. Li, “The study of APT attack stage model,” in 2016 IEEE/ACIS 15th International Conference on Computer and Information Science (ICIS), June 2016, pp. 1–5. [Online]. Available: <https://doi.org/10.1109/ICIS.2016.7550947>

Image Credits

["File:Celowanie do tarczy 2010 ubt.JPG"](#) by [Tsca.bot](#) is licensed under [CC BY 3.0](#)

["Circuit Board"](#) by [johnmuk](#) is licensed under [CC BY-NC-SA 2.0](#)

["Unstructured LAN Patching"](#) by [webernetz](#) is licensed under [CC BY 2.0](#)

["File:Quercus robur - sprouting acorn.jpg"](#) by [Amphis](#) is licensed under [CC BY-SA 3.0](#)

["File:Magnifying glass.jpg"](#) by [Tomomarusan](#) is licensed under [CC BY 2.5](#)

["File:Rockdale County Jail detail of jail entrance at SE corner.jpg"](#) by [Krelnik](#) is licensed under [CC BY-SA 4.0](#)

["DSC_0093"](#) by [bobosh_t](#) is licensed under [CC BY-NC 2.0](#)

["File:Phone booth, Lidgett Park Road, Leeds - DSC07610.JPG"](#) by [Green Lane](#) is licensed under [CC BY-SA 3.0](#)

["File:Snort.jpg"](#) by [Tegshbuyan](#) is licensed under [CC BY-SA 4.0](#)

Image Credits

"[Wazuh_blue.png](#)" by Zenidd is licensed under [CC BY-SA 4.0](#)

"[Gears, at the Musée des Arts et Métiers](#)" by [ArkanGL](#) is licensed under [CC BY-NC-SA 2.0](#)

"[wp plugins](#)" by [Sean MacEntee](#) is licensed under [CC BY 2.0](#)

"[Basset Hound Crooning](#)" by [jeffmgrandy](#) is licensed under [CC BY-NC-ND 2.0](#)

"[Random Blue Sky with White Clouds](#)" by [warpdesign](#) is licensed under [CC BY 2.0](#)

"[File:2010-03-18 \(27\) Honey bee, Honigbiene, Apis mellifica.JPG](#)" by [VBuhl](#) is licensed under [CC BY-SA 3.0](#)

"[File:Oszilloscop.png](#)" by [Kae-ru](#) is licensed under [CC BY-SA 3.0](#)

"[Circles](#)" by [places_lost](#) is licensed under [CC BY 2.0](#)

"[Index](#)" by [Ben Weiner](#) is licensed under [CC BY-ND 2.0](#)

Image Credits

"[IMG_4813](#)" by [bjcoving](#) is licensed under [CC BY-NC-ND 2.0](#)

"[Stethoscope](#)" by [surroundsound5000](#) is licensed under [CC BY-NC 2.0](#)

"[Money!](#)" by [bmJi](#) is licensed under [CC BY-NC-ND 2.0](#)

"[mind the gap](#)" by [osde8info](#) is licensed under [CC BY-SA 2.0](#)